

bfs_omp.cpp 中的 bfs_omp() 函数及其调用部分

```
1 #include "bfs_common.h"
2 #include "graph.h"
3 #include <stdio>
4 #include <omp.h>
5
6 #define ROOT_NODE_ID 0
7 #define NOT_VISITED_MARKER -1
8
9 void my_vertex_set_clear(vertex_set *list) { list->count = 0; }
10
11 void my_vertex_set_init(vertex_set *list, int count) {
12     list->max_vertices = count;
13     list->vertices = (int *)malloc(sizeof(int) * list->max_vertices);
14     my_vertex_set_clear(list);
15 }
16
17 // Take one step of "top-down" BFS. For each vertex on the frontier,
18 // follow all outgoing edges, and add all neighboring vertices to the
19 // new_frontier.
20 void my_top_down_step(Graph g, vertex_set *frontier, vertex_set *new_frontier,
21                       int *distances) {
22
23     for (int i = 0; i < frontier->count; i++) {
24
25         int node = frontier->vertices[i];
26
27         int start_edge = g->outgoing_starts[node];
28         int end_edge = (node == g->num_nodes - 1) ? g->num_edges
29                                                         : g->outgoing_starts[node + 1];
30
31         // attempt to add all neighbors to the new frontier
32         for (int neighbor = start_edge; neighbor < end_edge; neighbor++) {
33             int outgoing = g->outgoing_edges[neighbor];
34
35             if (distances[outgoing] == NOT_VISITED_MARKER) {
36                 distances[outgoing] = distances[node] + 1;
37                 int index = new_frontier->count++;
38                 new_frontier->vertices[index] = outgoing;
39             }
40         }
41     }
```

```

41     }
42 }
43
44 void my_bottom_up_step(Graph g, vertex_set *new_frontier,
45                       int *distances, int num_threads, vertex_set *list,
46                       double beta, int distance) {
47     #pragma omp parallel for schedule(guided)
48     for (int i = 0; i < num_threads; ++i) {
49         my_vertex_set_clear(list + i);
50     }
51     #pragma omp parallel for schedule(guided)
52     for (int i = 0; i < num_nodes(g); ++i) {
53         if (distances[i] != NOT_VISITED_MARKER) continue;
54         int end = (i == g->num_nodes - 1) ? g->num_edges : g->incoming_starts[i + 1];
55         for (int v = g->incoming_starts[i]; v < end; ++v) {
56             if (distances[g->incoming_edges[v]] == distance) {
57                 int id = omp_get_thread_num();
58                 distances[i] = distance + 1;
59                 list[id].vertices[list[id].count++] = i;
60                 break;
61             }
62         }
63     }
64     int sum = 0;
65     #pragma omp parallel for schedule(guided) reduction(+:sum)
66     for (int i = 0; i < num_threads; ++i)
67         sum += list[i].count;
68     if (sum > 1. * num_nodes(g) / beta) {
69         new_frontier->count = sum;
70         return;
71     }
72     for (int i = 0; i < num_threads; ++i) {
73         for (int j = 0; j < list[i].count; ++j) {
74             new_frontier->vertices[new_frontier->count++] = list[i].vertices[j];
75         }
76     }
77 }
78
79 void bfs_omp(Graph graph, solution *sol) {
80     /** Your code ... */
81     int num_threads;
82     #pragma omp parallel
83     {
84         #pragma omp master
85         num_threads = omp_get_num_threads();
86     }
87     double beta = 120;
88     vertex_set list1;

```

```

89     vertex_set list2;
90     vertex_set *list = new vertex_set[num_threads];
91     my_vertex_set_init(&list1, graph->num_nodes);
92     my_vertex_set_init(&list2, graph->num_nodes);
93     #pragma omp parallel for schedule(guided)
94     for (int i = 0; i < num_threads; ++i)
95         my_vertex_set_init(list + i, graph->num_nodes);
96
97     vertex_set *frontier = &list1;
98     vertex_set *new_frontier = &list2;
99
100    // initialize all nodes to NOT_VISITED
101    #pragma omp parallel for schedule(guided)
102    for (int i = 0; i < graph->num_nodes; i++)
103        sol->distances[i] = NOT_VISITED_MARKER;
104
105    // setup frontier with the root node
106    frontier->vertices[frontier->count++] = ROOT_NODE_ID;
107    sol->distances[ROOT_NODE_ID] = 0;
108
109    int distance = 0;
110
111    while (frontier->count != 0) {
112        my_vertex_set_clear(new_frontier);
113        if (frontier->count > 1. * num_nodes(graph) / beta)
114            my_bottom_up_step(graph, new_frontier, sol->distances, num_threads, list,
115            beta, distance);
116        else
117            my_top_down_step(graph, frontier, new_frontier, sol->distances);
118
119        // swap pointers
120        vertex_set *tmp = frontier;
121        frontier = new_frontier;
122        new_frontier = tmp;
123        ++distance;
124    }
125 }

```

实现思路为：

1. 首先实现 Bottom Up 方式的 BFS，即枚举所有结点，看是否是当前 frontier 集合中结点的邻居并且没有访问。
2. 对于结点的枚举采用 OpenMP 加速，即对于最外层循环使用 #pragma omp parallel for，而在寻找到可行的结点时，加入线程对应编号的 list 中以减少同步；而由于不同的不同的结点分配给了不同的 OpenMP 线程，因此不会出现 distances 和 new_frontier 的写入竞争。
3. 设置 beta 参数，当 frontier 中的节点数超过 $\frac{n}{\beta}$ 时，再采用 Bottom Up 方式加速，否则使用单线程的 Top Down 方式进行朴素计算。

bfs_omp_mpi.cpp 中的 *bfs_omp_mpi()* 函数及其调用部分

```
1  #include "bfs_common.h"
2  #include "graph.h"
3  #include <stdio>
4  #include <sys/time.h>
5  #include <mpi.h>
6  #include <omp.h>
7
8  #define ROOT_NODE_ID 0
9  #define NOT_VISITED_MARKER -1
10
11 void my_vertex_set_clear(vertex_set *list) { list->count = 0; }
12
13 void my_vertex_set_init(vertex_set *list, int count) {
14     list->max_vertices = count;
15     list->vertices = (int *)malloc(sizeof(int) * list->max_vertices);
16     my_vertex_set_clear(list);
17 }
18
19 // Take one step of "top-down" BFS. For each vertex on the frontier,
20 // follow all outgoing edges, and add all neighboring vertices to the
21 // new_frontier.
22 void my_top_down_step(Graph g, vertex_set *frontier, vertex_set *new_frontier,
23                     int *distances, int rank, int beta) {
24     int distance = distances[frontier->vertices[0]];
25
26     if (rank == 0) {
27         for (int i = 0; i < frontier->count; i++) {
28
29             int node = frontier->vertices[i];
30
31             int start_edge = g->outgoing_starts[node];
32             int end_edge = (node == g->num_nodes - 1) ? g->num_edges
33                     : g->outgoing_starts[node + 1];
34
35             // attempt to add all neighbors to the new frontier
36             for (int neighbor = start_edge; neighbor < end_edge; neighbor++) {
37                 int outgoing = g->outgoing_edges[neighbor];
38                 if (distances[outgoing] == NOT_VISITED_MARKER) {
39                     distances[outgoing] = distances[node] + 1;
40                     int index = new_frontier->count++;
41                     new_frontier->vertices[index] = outgoing;
42                 }
43             }
44         }
45     }
```

```

46 MPI_Bcast(&new_frontier->count, 1, MPI_INT, 0, MPI_COMM_WORLD);
47 if (new_frontier->count <= 1. * num_nodes(g) / beta) {
48     MPI_Bcast(new_frontier->vertices, new_frontier->count, MPI_INT, 0,
MPI_COMM_WORLD);
49     #pragma omp parallel for schedule(guided)
50     for (int i = 0; i < new_frontier->count; ++i)
51         distances[new_frontier->vertices[i]] = distance + 1;
52 } else {
53     MPI_Bcast(new_frontier->vertices, new_frontier->count, MPI_INT, 0,
MPI_COMM_WORLD);
54     #pragma omp parallel for schedule(guided)
55     for (int i = 0; i < new_frontier->count; ++i)
56         distances[new_frontier->vertices[i]] = distance + 1;
57     // MPI_Bcast(distances, num_nodes(g), MPI_INT, 0, MPI_COMM_WORLD);
58 }
59 }
60
61 void my_bottom_up_step(Graph g, vertex_set *new_frontier, vertex_set *tmp_frontier,
62     int *distances, int num_threads, vertex_set *list,
63     int rank, int nprocs, int *count,
64     int *disp, double beta, int distance) {
65     // timeval start, end;
66     // gettimeofday(&start, NULL);
67
68     int blockSize = num_nodes(g) / nprocs;
69     if (num_nodes(g) % nprocs) ++blockSize;
70     int l = std::min(num_nodes(g), blockSize * rank), r = std::min(num_nodes(g) - 1,
blockSize * (rank + 1));
71     #pragma omp parallel for schedule(guided)
72     for (int i = 0; i < num_threads; ++i) {
73         my_vertex_set_clear(list + i);
74     }
75     #pragma omp parallel for schedule(guided)
76     for (int i = l; i < r; ++i) {
77         if (distances[i] != NOT_VISITED_MARKER) continue;
78
79         for (int *v = g->incoming_edges + g->incoming_starts[i]; v < g->incoming_edges
+ g->incoming_starts[i + 1]; ++v) {
80             if (distances[*v] == distance) {
81                 int id = omp_get_thread_num();
82                 distances[i] = distance + 1;
83                 list[id].vertices[list[id].count++] = i;
84                 break;
85             }
86         }
87     }
88     if ((num_nodes(g) - 1) / blockSize == rank && distances[num_nodes(g) - 1] ==
NOT_VISITED_MARKER) {

```

```

89     int i = num_nodes(g) - 1;
90     for (int v = g->incoming_starts[i]; v < g->num_edges; ++v) {
91         if (distances[g->incoming_edges[v]] == distance) {
92             distances[i] = distance + 1;
93             list[0].vertices[list[0].count++] = i;
94             break;
95         }
96     }
97 }
98 int sum = 0;
99 #pragma omp parallel for schedule(guided) reduction(+:sum)
100 for (int i = 0; i < num_threads; ++i)
101     sum += list[i].count;
102 // gettimeofday(&end, NULL);
103 // printf("Full Rank %d: %lf ms.\n",
104 //         rank, (1000000.0 * (end.tv_sec - start.tv_sec) + end.tv_usec -
start.tv_usec) / 1000.0);
105 MPI_Allreduce(&sum, &sum, 1, MPI_INT, MPI_SUM, MPI_COMM_WORLD);
106 if (sum > 1. * num_nodes(g) / beta) {
107     new_frontier->count = sum;
108     MPI_Allgather(distances + 1, blockSize, MPI_INT, distances, blockSize, MPI_INT,
MPI_COMM_WORLD);
109     return;
110 }
111 my_vertex_set_clear(tmp_frontier);
112 for (int i = 0; i < num_threads; ++i) {
113     for (int j = 0; j < list[i].count; ++j) {
114         tmp_frontier->vertices[tmp_frontier->count++] = list[i].vertices[j];
115     }
116 }
117 MPI_Allgather(&tmp_frontier->count, 1, MPI_INT, count, 1, MPI_INT,
MPI_COMM_WORLD);
118 for (int i = 0; i < nprocs; ++i) {
119     if (i) disp[i] = disp[i - 1] + count[i - 1];
120     else disp[i] = 0;
121 }
122 MPI_Allgatherv(tmp_frontier->vertices, tmp_frontier->count, MPI_INT,
new_frontier->vertices, count, disp, MPI_INT, MPI_COMM_WORLD);
123 new_frontier->count = disp[nprocs - 1] + count[nprocs - 1];
124 #pragma omp parallel for schedule(guided)
125 for (int i = 0; i < new_frontier->count; ++i)
126     distances[new_frontier->vertices[i]] = distance + 1;
127 }
128
129 void bfs_omp_mpi(Graph graph, solution *sol) {
130     /** Your code ... */
131     int rank, nprocs;
132     MPI_Comm_rank(MPI_COMM_WORLD, &rank);

```

```

133 MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
134
135 int blockSize = num_nodes(graph) / nprocs;
136 if (num_nodes(graph) % nprocs) ++blockSize;
137
138 int *count = new int[nprocs];
139 int *disp = new int[nprocs];
140 int num_threads;
141 #pragma omp parallel
142 {
143     #pragma omp master
144     num_threads = omp_get_num_threads();
145 }
146 double beta = 120;
147 vertex_set list1;
148 vertex_set list2;
149 vertex_set list3;
150 vertex_set *list = new vertex_set[num_threads];
151 my_vertex_set_init(&list1, graph->num_nodes);
152 my_vertex_set_init(&list2, graph->num_nodes);
153 my_vertex_set_init(&list3, blockSize);
154 #pragma omp parallel for schedule(guided)
155 for (int i = 0; i < num_threads; ++i)
156     my_vertex_set_init(list + i, graph->num_nodes);
157
158 vertex_set *frontier = &list1;
159 vertex_set *new_frontier = &list2;
160 vertex_set *tmp_frontier = &list3;
161
162 int *my_distances = new int[nprocs * blockSize];
163 // initialize all nodes to NOT_VISITED
164 #pragma omp parallel for schedule(guided)
165 for (int i = 0; i < graph->num_nodes; i++)
166     my_distances[i] = NOT_VISITED_MARKER;
167
168 // setup frontier with the root node
169 frontier->vertices[frontier->count++] = ROOT_NODE_ID;
170 my_distances[ROOT_NODE_ID] = 0;
171
172 int distance = 0;
173
174 // printf("Rank %d Start!!!!\n", rank);
175
176 while (frontier->count != 0) {
177     my_vertex_set_clear(new_frontier);
178
179     // timeval start, end;
180     // gettimeofday(&start, NULL);

```

```

181
182     if (frontier->count > 1. * num_nodes(graph) / beta)
183         // printf("Bottom up "),
184         my_bottom_up_step(graph, new_frontier, tmp_frontier, my_distances,
num_threads, list, rank, nprocs, count, disp, beta, distance);
185     else
186         // printf("Top down "),
187         my_top_down_step(graph, frontier, new_frontier, my_distances, rank, beta);
188
189     // gettimeofday(&end, NULL);
190     // printf("Rank %d: %lf ms. %d\n",
191     //         rank, (1000000.0 * (end.tv_sec - start.tv_sec) + end.tv_usec -
start.tv_usec) / 1000.0, new_frontier->count);
192
193     // swap pointers
194     vertex_set *tmp = frontier;
195     frontier = new_frontier;
196     new_frontier = tmp;
197     ++distance;
198 }
199 #pragma omp parallel for schedule(guided)
200 for (int i = 0; i < graph->num_nodes; i++)
201     sol->distances[i] = my_distances[i];
202 // printf("Rank %d End!!!!\n", rank);
203 }
204

```

实现思路为：

1. 与 bfs_omp() 大体一样，但在 Bottom Up 的实现中，预分配给每个进程固定的结点编号序列用以更新。
2. 对于每个进程，仍然采用 OpenMP 加速并合并。
3. 在 Bottom Up 方式结束时，通过 MPI_Allgather() 和 MPI_Allgatherv() 进行通讯，更新本次迭代的 new_frontier 以及 distances 数组。
4. 在 Top Down 方式结束时，通过 MPI_Bcast() 进行对应信息的广播。

性能优化与效果

1. 通过设置阈值 β 来对于不同情况选择 Top Down 与 Bottom Up 两种方式。在 OpenMP 版本的测试中，可以获得 30 ~ 40% 左右的加速。
2. 通过调整 β 与 num_threads 的数值。可以在 OpenMP+MPI 版本中获得 20% 左右的加速。
3. 在 Bottom Up 方式结束前，首先判断是否 new_frontier 的总点数是否达到了设定的阈值要求即 $\frac{n}{\beta}$ ，如果达到则不进行线程/进程之间 new_frontier 数组的合并，而只是记录总点数（MPI 通讯时通过 MPI_Allgather() 同步 distances 数组）。在 graph/68m.graph 数据集下，两种实现大致均可获得 10ms 左右加速。

OpenMP 不同线程数运行结果

在 graph/200m.graph 测试集下，采用 28 线程数，耗时为 297.3816ms。

注：下面测试中时间均为 graph/500m.graph 用时。

线程数	运行时间(ms)	加速比
1	10595.4758	1.00
7	2349.3291	4.51
14	1244.3149	8.52
28	774.3301	13.68

OpenMP+MPI 不同进程数运行结果

在 graph/500m.graph 测试集下，采用 4×1 进程数，耗时为 360.2648ms。

在 graph/200m.graph 测试集下，采用 4×1 进程数，耗时为 239.1775ms。

注：下面测试中时间均为 graph/68m.graph 用时。

机器数 × 进程数	运行时间(<i>ms</i>)	加速比
1 × 1	46.4430	1.00
1 × 2	45.4541	1.02
1 × 4	77.4105	0.60
1 × 14	95.2084	0.49
1 × 28	121.5943	0.38
2 × 1	39.3658	1.18
2 × 2	52.8389	0.88
2 × 4	57.9481	0.80
2 × 14	89.0764	0.52
2 × 28	126.0286	0.37
4 × 1	32.8034	1.42
4 × 2	44.6022	1.04
4 × 4	57.0379	0.81
4 × 14	94.0943	0.49
4 × 28	322.9891	0.14